
CS 224N: Assignment 1

RYAN McMAHON

SATURDAY 28TH JANUARY, 2017

Contents

1	Problem 1: Softmax (10 pts)	2
1.1	(a) Softmax Invariance to Constant (5 pts)	2
1.2	(b) Softmax Coding (5 pts)	2
2	Problem 2: Neural Network Basics (30 pts)	3
2.1	(a) Sigmoid Gradient (3 pts)	3
2.2	(b) Softmax Gradient with Cross Entropy Loss(3 pts)	4

Problem 1: Softmax (10 pts)

1.1 (a) Softmax Invariance to Constant (5 pts)

Prove that softmax is invariant to constant offsets in the input, that is, for any input vector x and any constant c , $\text{softmax}(x) = \text{softmax}(x + c)$, where $x + c$ means adding the constant c to every dimension of x . Remember that

$$\text{softmax}(x)_i = \frac{e^{x_i}}{\sum_j e^{x_j}} \quad (1.1)$$

Answer:

We can show that $\text{softmax}(x) = \text{softmax}(x + c)$ by factoring out c and canceling:

$$\begin{aligned} \text{softmax}(x + c)_i &= \frac{e^{x_i + c}}{\sum_j e^{x_j + c}} = \frac{e^{x_i} \times e^c}{e^c \times \sum_j e^{x_j}} \\ &= \frac{e^{x_i} \times \cancel{e^c}}{\cancel{e^c} \times \sum_j e^{x_j}} = \text{softmax}(x)_i \end{aligned}$$

1.2 (b) Softmax Coding (5 pts)

Given an input matrix of N rows and D columns, compute the softmax prediction for each row using the optimization in part (a). Write your implementation in `q1_softmax.py`. You may test by executing `python q1_softmax.py`.

Note: The provided tests are not exhaustive. Later parts of the assignment will reference this code so it is important to have a correct implementation. Your implementation should also be efficient and vectorized whenever possible (i.e., use numpy matrix operations rather than for loops). A non-vectorized implementation will not receive full credit!

Answer:

See code: `~/code/q1_softmax.py`.

Problem 2: Neural Network Basics (30 pts)

2.1 (a) Sigmoid Gradient (3 pts)

Derive the gradients of the sigmoid function and show that it can be rewritten as a function of the function value (i.e., in some expression where only $\sigma(x)$, but not x , is present). Assume that the input x is a scalar for this question. Recall, the sigmoid function is

$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad (2.1)$$

Answer:

$$\begin{aligned} \sigma(x) &= \frac{1}{1 + e^{-x}} \\ &= \frac{e^x}{1 + e^x} \\ \frac{\partial}{\partial x} \sigma(x) &= \frac{e^x \times (1 + e^x) - (e^x \times e^x)}{(1 + e^x)^2} \\ &= \frac{e^x + \cancel{(e^x \times e^x)} - \cancel{(e^x \times e^x)}}{(1 + e^x)^2} \\ &= \frac{e^x}{(1 + e^x)^2} = \sigma(x) \times (1 - \sigma(x)) \end{aligned}$$

Because $1 - \sigma(x) = \sigma(-x)$ we can show that:

$$\begin{aligned} \frac{\partial}{\partial x} \sigma(x) &= \frac{e^x}{(1 + e^x)^2} \\ &= \sigma(x) \times \sigma(-x) \\ &= \frac{e^x}{1 + e^x} \times \frac{1}{1 + e^{+x}} \\ &= \frac{e^x}{(1 + e^x)^2} \end{aligned}$$

2.2 (b) Softmax Gradient w/ Cross Entropy Loss (3 pts)

Derive the gradient with regard to the inputs of a softmax function when cross entropy loss is used for evaluation, i.e., find the gradients with respect to the softmax input vector θ , when the prediction is made by $\hat{\mathbf{y}} = \text{softmax}(\theta)$. Remember the cross entropy function is

$$CE(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_i y_i \times \log(\hat{y}_i) \quad (2.2)$$

where \mathbf{y} is the one-hot label vector; and $\hat{\mathbf{y}}$ is the predicted probability vector for all classes. (Hint: you might want to consider the fact many elements of \mathbf{y} are zeros, and assume that only the k -th dimension of \mathbf{y} is one.)