CS 224N: Assignment 1

RYAN MCMAHON   SUNDAY 29TH JANUARY, 2017

Contents

Problem 1: Softmax (10 pts)

1.1 (a) Softmax Invariance to Constant (5 pts)

Prove that softmax is invariant to constant offsets in the input, that is, for any input vector x and any constant c, softmax(x) = softmax($x + c$), where $x + c$ means adding the constant c to every dimension of x. Remember that

$$softmax(x)_i = \frac{e^{x_i}}{\sum_j e^{x_j}} \tag{1.1}$$

Answer:

We can show that softmax(x) = softmax($x + c$) by factoring out c and canceling:

$$softmax(x + c)_i = \frac{e^{x_i + c}}{\sum_j e^{x_j + c}} = \frac{e^{x_i} \times e^c}{e^c \times \sum_j e^{x_j}}$$

$$= \frac{e^{x_i} \times \cancel{e^c}}{\cancel{e^c} \times \sum_j e^{x_j}} = softmax(x)_i$$

1.2 (b) Softmax Coding (5 pts)

Given an input matrix of N rows and D columns, compute the softmax prediction for each row using the optimization in part (a). Write your implementation in `q1_softmax.py`*. You may test by executing* `python q1_softmax.py`*.*

Note: The provided tests are not exhaustive. Later parts of the assignment will reference this code so it is important to have a correct implementation. Your implementation should also be efficient and vectorized whenever possible (i.e., use numpy matrix operations rather than for loops). A non-vectorized implementation will not receive full credit!

Answer:

See code: $\sim$/code/q1_softmax.py.

Problem 2: Neural Network Basics (30 pts)

2.1 (a) Sigmoid Gradient (3 pts)

Derive the gradients of the sigmoid function and show that it can be rewritten as a function of the function value (i.e., in some expression where only (x), but not x, is present). Assume that the input x is a scalar for this question. Recall, the sigmoid function is

$$\sigma(x) = \frac{1}{1 + e^{-x}} \tag{2.1}$$

Answer:

$$
\begin{aligned}
\sigma(x) &= \frac{1}{1 + e^{-x}} \\
&= \frac{e^x}{1 + e^x} \\
\frac{\partial}{\partial x}\sigma(x) &= \frac{e^x \times (1 + e^x) - (e^x \times e^x)}{(1 + e^x)^2} \\
&= \frac{e^x + \cancel{(e^x \times e^x)} - \cancel{(e^x \times e^x)}}{(1 + e^x)^2} \\
&= \frac{e^x}{(1 + e^x)^2} = \sigma(x) \times (1 - \sigma(x))
\end{aligned}
$$

Because $1 - \sigma(x) = \sigma(-x)$ we can show that:

$$
\begin{aligned}
\frac{\partial}{\partial x}\sigma(x) &= \frac{e^x}{(1 + e^x)^2} \\
&= \sigma(x) \times \sigma(-x) \\
&= \frac{e^x}{1 + e^x} \times \frac{1}{1 + e^{+x}} \\
&= \frac{e^x}{(1 + e^x)^2}
\end{aligned}
$$

2.2 (b) Softmax Gradient w/ Cross Entropy Loss (3 pts)

Derive the gradient with regard to the inputs of a softmax function when cross entropy loss is used for evaluation, i.e., find the gradients with respect to the softmax input vector $\boldsymbol{\theta}$, when the prediction is made by $\hat{\mathbf{y}} = softmax(\boldsymbol{\theta})$. Remember the cross entropy function is

$$CE(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_i y_i \times log(\hat{y}_i) \tag{2.2}$$

where $\mathbf{y}$ is the one-hot label vector, and $\hat{y}$ is the predicted probability vector for all classes. (Hint: you might want to consider the fact many elements of $\mathbf{y}$ are zeros, and assume that only the $k-th$ dimension of $\mathbf{y}$ is one.)

Answer:

Let S represent the softmax function:

$$f_i = e^{\theta_i}$$

$$g_i = \sum_{k=1}^{K} e^{\theta_k}$$

$$S_i = \frac{f_i}{g_i}$$

$$\frac{\partial S_i}{\partial \theta_j} = \frac{f_i' g_i - g_i' f_i}{g_i^2}$$

So if $i = j$:

$$f_i' = f_i; \ g_i' = e^{\theta_j}$$

$$\frac{\partial S_i}{\partial \theta_j} = \frac{e^{\theta_i} \sum_k e^{\theta_k} - e^{\theta_j} e^{\theta_i}}{(\sum_k e^{\theta_k})^2}$$

$$= \frac{e^{\theta_i}}{\sum_k e^{\theta_k}} \times \frac{\sum_k e^{\theta_k} - e^{\theta_j}}{\sum_k e^{\theta_k}}$$

$$= S_i \times (1 - S_i)$$

And if $i \neq j$:

$$\frac{\partial S_i}{\partial \theta_j} = \frac{0 - e^{\theta_j} e^{\theta_i}}{(\sum_k e^{\theta_k})^2}$$

$$= -\frac{e^{\theta_j}}{\sum_k e^{\theta_k}} \times \frac{e^{\theta_i}}{\sum_k e^{\theta_k}}$$

$$= -S_j \times S_i$$

We can now use these when operating on our loss function (let L represent the cross entropy function):

$$\frac{\partial L}{\partial \theta_i} = -\sum_k y_k \frac{\partial log S_k}{\partial \theta_i}$$

$$= -\sum_k y_k \frac{1}{S_k} \frac{\partial S_k}{\partial \theta_i}$$

$$= -y_i(1 - S_i) - \sum_{k \neq i} y_k \frac{1}{S_k}(-S_k \times S_i)$$

$$= -y_i(1 - S_i) + \sum_{k \neq i} y_k S_i$$

$$= -y_i + y_i S_i + \sum_{k \neq i} y_k S_i$$

$$= S_i(\sum_k y_k) - y_i$$

And because we know that $\sum_k y_k = 1$:

$$\frac{\partial L}{\partial \theta_i} = S_i - y_i$$

2.3 (c) One Hidden Layer Gradient (6 pts)

Derive the gradients with respect to the inputs x to a one-hidden-layer neural network (that is, find $\frac{\partial J}{\partial x}$ where $J = CE(\mathbf{y}, \hat{\mathbf{y}})$ is the cost function for the neural network). The neural network employs sigmoid activation function for the hidden layer, and softmax for the output layer. Assume the one-hot label vector is $\mathbf{y}$, and cross entropy cost is used. (Feel free to use $\sigma'(x)$ as the shorthand for sigmoid gradient, and feel free to define any variables whenever you see fit.)

Recall that forward propoagation is as follows

$$\mathbf{h} = sigmoid(\boldsymbol{x}\boldsymbol{W}_1 + \boldsymbol{b}_1) \qquad\qquad \hat{\boldsymbol{y}} = softmax(\boldsymbol{h}\boldsymbol{W}_2 + \boldsymbol{b}_2)$$

Answer:

Let $f_2 = \boldsymbol{x}\boldsymbol{W}_1 + \boldsymbol{b}_1$ and $f_3 = \boldsymbol{h}\boldsymbol{W}_2 + \boldsymbol{b}_2$;

$$\frac{\partial J}{\partial f_3} = \delta_3 = \hat{\boldsymbol{y}} - \boldsymbol{y}$$

$$\frac{\partial J}{\partial f_2} = \delta_2 = \boldsymbol{W}_2\delta_3 \circ \sigma'(f_2)$$

$$\frac{\partial J}{\partial \boldsymbol{x}} = \frac{\partial J}{\partial f_2}\frac{\partial f_2}{\partial \boldsymbol{x}}$$

$$= \delta_2\frac{\partial f_2}{\partial \boldsymbol{x}}$$

$$= \delta_2\frac{\partial}{\partial \boldsymbol{x}}\boldsymbol{x}\boldsymbol{W}_1 + \boldsymbol{b}_1$$

$$= \delta_2\boldsymbol{W}_1$$

2.4 (d) No. Parameters (2 pts)

*How many parameters are there in this neural network [from (**c**) above], assuming the input is D_x-dimensional, the output is D_y-dimensional, and there are H hidden units?*

Answer:

$$n_{W_1} = D_x \times H$$

$$n_{b_1} = H$$

$$n_{W_2} = H \times D_y$$

$$n_{b_2} = D_y$$

$$N = (D_x \times H) + H + (H \times D_y) + D_y$$

2.5 (e) Sigmoid Activation Code (4 pts)

Fill in the implementation for the sigmoid activation function and its gradient in `q2_sigmoid.py`. *Test your implementation using* `python q2_sigmoid.py`. *Again, thoroughly test your code as the provided tests may not be exhaustive.*

Answer:

See code: ~/code/q2_sigmoid.py.

2.6 (f) Gradient Check Code (4 pts)

To make debugging easier, we will now implement a gradient checker. Fill in the implementation for `gradcheck_naive` *in* `q2_gradcheck.py`. *Test your code using* `python q2_gradcheck.py`.

Answer:

See code: ~/code/q2_gradcheck.py.